

EE122 Project 3a

(Version 1.2)

Due: 5pm, November 18th, 2013 (hard deadline)

Sangjin Han Steve Wang Kaifei Chen Aurojit Panda Jamie Tsao

Overview

In this project, you will implement a basic firewall running at end hosts. A firewall is a “*security system that controls the incoming and outgoing network traffic by analyzing the data packets and determining whether they should be allowed through or not, based on a rule set*” [Wikipedia]. Unlike the previous projects of this course, where you worked in simulated environments, you will deal with real packets in a Linux-based virtual machine (VM) for this project.

This month-long project is divided into two parts, and each part has its own submission deadline. In the first part (3a), which is covered in this document, you are asked to build a stateless firewall on top of the given framework. In the second part (3b), you will be extending the functionality of your firewall to support stateful rules at the application layer. Note that your solution for Project 3a will also be used as a base for Project 3b. It is very important to keep your code readable and extensible.

Your task for Project 3a is to implement a firewall that filters out packets based on simple firewall (Protocol/IP/Port and DNS query) rules on a packet-by-packet basis. Upon successful completion of this project, you will be able to:

- Understand the basic functionalities of a firewall.
- Be familiar with the details of TCP/IP packet formats.
- Explore low-level packet processing.
- Utilize various tools for network testing.

Besides writing code, you will need to (and should) spend a lot of time to understand protocol specifications, to design algorithms, and to test your application. Start working on the project as soon as possible.

Changelog

Always check the latest version of this document. It is your responsibility that your solution conforms to the latest version of the spec.

v1.2 (11/09/2013)

- Clarifications:
 - DNS rules may appear before Protocol/IP/Port rules.
 - If a DNS query packet is too malformed to decode, you should drop the packet, even if there is no DNS rules.
 - DNS query packets may have entries in non-question sections.

v1.1 (11/06/2013)

- Fixed typos:
 - Use “`sudo ./main.py`” instead of “`sudo main.py`”
- Clarifications:
 - The rules file is not fixed and can have arbitrary rules.
 - The performance measurement will also consider the startup time of your solution.
 - We may use a different `geoipdb.txt` file from the provided one.
 - Your program should not crash, even with malformed packets.
- Added `dig` in the References section
- Performance measurement interval: “> 10 seconds” → “> 20 seconds”
- Added a link to the supplementary document:
<http://www-inst.eecs.berkeley.edu/~ee122/fa13/projects/project3a/supp.pdf>

v1.0 (10/30/2013)


- First release

VM Setup

A personal firewall must have the ability to intercept incoming and outgoing network packets from network interfaces. Since this operation is security-critical and dependent on your operating system, we provide a VM image that is preconfigured to be readily used for the project. To run the VM, you will need to use your personal x86 laptop/desktop, rather than instructional machines. The VM runs on the Ubuntu Linux 13.04 Desktop edition.

Understanding the VM Network Configuration

The following figure illustrates how the VM network is configured.



The arrows in the diagram represent the flow of network traffic (packets). In the VM, there are two network interfaces, namely `ext` and `int` (short for “external” and “internal”, respectively), and the firewall in between, as a bump-in-the-wire. All outgoing packets of the VM will be sent through the `int` interface. Similarly, all incoming packets will be received via the `ext` interface. Your firewall should inspect packets received from one interface and selectively pass them to the other interface. For example, when the firewall receives an incoming packet from `ext`, it determines whether to pass the packet through based on the firewall rules. If so, the firewall transmits the packet onto `int`, to be processed by the Linux TCP/IP stack. Similar things happen for outgoing packets as well.

When the firewall is not running, nothing relays packets between `ext` and `int`. **Thus the VM has no access to the outside network by default.** Don’t be surprised if you cannot access to any websites with Firefox in the VM. Once you have correctly implemented your firewall, or after running “`sudo ./main.py --mode bypass`” (explained later), your applications in the VM will have access to the Internet.

Note: Due to some internal issues, **you will not see any packets on the `int` interface with `tcpdump`/wireshark, while the firewall is not running.**

For simplicity, we made the following design decisions for the project:

- `ext` and `int` are Ethernet interfaces, which carry Ethernet frames (with 14-byte Ethernet header followed by its payload, mostly an IP packet). However, your firewall will only need to care about IP packets; it receives and sends IP packets, and all Ethernet-related operations will be handled by the code we have provided.
- IPv6 was completely disabled in the VM. All you see in the firewall is IPv4 packets.
- Your firewall will not receive any fragmented IP packets.

Another thing worth mentioning is that there is a NAT¹ module in VirtualBox, which connects the VM to the outside network. While the details of NAT is out of the scope of this project, there are two things to remember:

- The IP address of the VM (statically set to 10.0.2.15) will not be seen from the outside network. This is because the NAT module translates it into the IP address of the host (your laptop) for every outgoing packet, and the other way around for every incoming packet.
- By default, the NAT module will not allow incoming connections destined for the VM (e.g., the NAT module will not forward incoming TCP SYN packets). In other words, you cannot run network server applications in the VM.

VM Installation

1. Download and install the latest version (**4.3 is recommended**) of VirtualBox, from <https://www.virtualbox.org/wiki/Downloads>
2. Download the VM image, from <http://bit.ly/ee122p3>
 - a. It may take a while to download the image, due to its large size (> 1.5 GB). We recommend you download on campus for faster speeds.
 - b. Make sure you have enough disk space, as the compressed image gets bigger when it is imported. We recommend securing at least 10 GB of free space.
3. Launch VirtualBox, and import the VM.
 - a. Select “File → Import Appliance”
 - b. Click the “Open appliance...” button
 - c. Choose “EE122proj3.ova”
4. Once imported, select “EE122proj3” on the left panel of the main window, and click the “Start” button to launch the VM.

¹ Network Address Translation: http://en.wikipedia.org/wiki/Network_address_translation

Account

The username is “ee122” and its password is “F1rewa!!” (number one instead of letter I and exclamation mark instead of letter L), without quotes. Feel free to change the password if your left hand fingers get tired.

Many applications (including the firewall itself) you will run on the VM need root permission, as they need to perform privileged network operations. Your account is granted sudo privileges without the password.

(Optional) Setting Host-Only Interface

You can skip this configuration if you do not mind working in the Ubuntu desktop environment.

Some of you may want to connect to the VM via SSH to work on the project, rather than directly using the GUI of the VM. For the connection between the host and the VM, you need to follow the procedure (based on Mac OS X, but it should be similar on other operating systems) below.

1. Turn off the VM.
2. Add a host-only interface for the host
 - a. Select “VirtualBox → Preferences” in the menu.
 - b. Choose the “Network” tab.
 - c. Choose the “Host-only Networks” tab.
 - d. Click the “+” button on the right.
 - e. Click the screwdriver button on the right.
 - f. On the “Adaptor” tab,
 - i. Set the “IPv4 Address” to 172.16.122.1
 - ii. Set the “IPv4 Network Mask” to 255.255.255.0
 - g. On the “DHCP” tab,
 - i. Unmark the “Enable Server” checkbox.
3. Add a host-only interface for the VM
 - a. Select the “EE122proj3” VM on the left panel of the VirtualBox Manager window.
 - b. Click the “Settings” button.
 - c. Choose the “Network” tab.
 - d. Choose the “Adaptor 2” tab.
 - e. Mark the “Enable Network Adaptor” checkbox.
 - i. Attached to: Host-only Adapter
 - ii. Name: vboxnet0 (or anything else you created above)
 - f. Click on the “Advanced”
 - g. Set the “MAC Address” to 080027107f8d (in hex)
 - h. Mark the “Cable Connected” checkbox.
4. Check if everything is OK.
 - a. Launch the VM

- b. Open a terminal window.
- c. Run the command `“ifconfig”`.
- d. You should be able to see the `“host”` interface.

The IP address of the host-only interface in the VM is 172.16.122.2. In the host (your laptop), you can make a SSH connection to the VM, with `“ssh ee122@172.16.122.2”` (on Linux or MAC) or your favorite SSH client (on Windows).

Network connections over this dedicated interface will not be affected by the firewall. All outgoing packets from the VM with a destination IP address in 172.16.122.0/24 will be sent through the `host` interface, instead of the `int` interface.

Project Specification

Provided Files

All files needed to do the project reside in the `/home/ee122/proj3` directory. Your task is to implement the project specification in the `firewall.py` file. Remember that this file is the only source code you submit for the project. All your modification must be done only in this file.

The `/home/ee122/backup_proj3` directory also contains the same files, just in case you need the original version.

`main.py`

This is the main executable file for your firewall. Modify this file only for debugging purposes. It contains some low-level code to intercept/inject packets from/to Linux network interfaces. Because of this, you will need root privileges to run this script.

```
sudo ./main.py [--mode <module name>] [--rule <rule file name>] [--opt1 arg1] ...
```

There are two predefined options:

- `--mode`: It specifies a Python module name that implements `Firewall` class. The default argument is “`firewall`”, which will run `firewall.py`.
- `--rule`: It specifies a rules file name. The default argument is “`rules.conf`”.

All other command-line options will be parsed and stored in the `config` dictionary, which is given to the constructor of `Firewall` class. For example, given “`--bwlimit 600k`” option (for Bells and Whistles 3), the value of `config[‘bwlimit’]` will be “`600k`”.

`bypass.py`

This is a dummy example that implements the bypass mode. In this mode, all incoming/outgoing packets will be passed regardless of the firewall rules. This mode can be useful when you want the VM to communicate without interference from the firewall, such as:

- Installing new applications
- Analyzing how network protocols work in normal conditions with `tcpdump`/`wireshark`
- Copying your source code to outside the VM for final submission
- ... so on.

The following command will run in the bypass mode:

```
sudo ./main.py --mode bypass
```

The `Firewall` class implemented in this file will give you some basic ideas on how to implement your own firewall in `firewall.py`, such as how to utilize the timer functionality or how to pass packets with the `.send_ip_packet()` method.

firewall.py

This is the file containing the skeleton for the code you need to implement. This module currently does not do anything, and all packets between the `int` and `ext` interfaces will be dropped. Your task is to complete the `Firewall` class in the file so that packets can be filtered out based on the firewall rules file. The skeleton of the class looks as follows.

```
class Firewall:
    def __init__(self, config, timer, iface_int, iface_ext):
        self.timer = timer
        self.iface_int = iface_int
        self.iface_ext = iface_ext
        ...

    def handle_timer(self):
        pass

    def handle_packet(self, pkt_dir, pkt):
        pass
```

`__init__()`:

- You should load the rules file (the filename is given in `config['rule']`).
- Also read `geoipdb.txt` here.

`handle_timer()`:

- When you schedule a timer with `self.timer.schedule(t)`, this handler will be invoked at time `t` (based on Python's `time.time()` function, which returns current time in seconds).
- You will not need to use the timer functionality for Project 3a, except for the Bells and Whistles 3.

`handle_packet()`:

- Whenever a packet is captured, this handler will be invoked.
- `pkt_dir` indicates the direction of the packet. It can be either of the following two values:
 - `PKT_DIR_INCOMING`: The packet has been received from the `ext` interface. You should call `self.iface_int.send_ip_packet()` to pass this packet.
 - `PKT_DIR_OUTGOING`: The packet has been received from the `int` interface. You should call `self.iface_ext.send_ip_packet()` to pass this packet.
- `pkt` is a Python string that contains the actual IP packet, including the IPv4 header.
- To drop the packet, simply omit the call to `.send_ip_packet()`.

Rules File Format

A firewall rule describes a type of packets that the firewall should pass/drop. A rules file contains firewall rules, each of which is written in a single line. `rules.conf` will be used by default, unless specified otherwise in the command line. The rules file decides whether to drop a packet or not, when the packet content/header matches one the defined rules. If multiple rules are matched, use the last one. There are two type of firewall rules: Protocol/IP/Port rules and DNS rules.

Note that the content of the rules file can be arbitrary. The provided `rules.conf` file in the `proj3` directory is just an example, and various rules files will be used to grade your solution. Your firewall should work correctly with any rule files that conform to the following format.

Protocol/IP/Port Rules

You apply Protocol/IP/Port rules for every packet that `Firewall.handle_packet()` takes. All protocol, external IP address, and external port fields must match to apply the verdict. “External” means “outside”, thus it may represent either the source or destination IP address/port, depending on the packet direction. For example, if an incoming packet has a UDP/IP header `8.8.4.4:53 → 10.0.2.15:32154`, the external IP is `8.8.4.4` and the external port is `53`.

Format: <verdict> <protocol> <external IP address> <external port>

Field	Possible values/formats	Description
verdict	1. “pass” 2. “drop”	<ul style="list-style-type: none">• <code>pass</code> means that a matched packet should be handed over to the interface on the other side, with the <code>send_ip_frame()</code> method (e.g., to <code>int</code> if received from <code>ext</code>).• <code>drop</code> discards the packet
protocol	1. “tcp” 2. “udp” 3. “icmp”	<ul style="list-style-type: none">• This field examines the Protocol field in the IPv4 header of packets.
external IP address	1. “any” 2. a 2-byte country code 3. a single IP address (e.g., <code>128.32.244.172</code>) 4. an IP prefix (e.g., <code>123.34.128.0/17</code>)	<ul style="list-style-type: none">• For country codes, see the “GeoIP DB” section.• <code>any</code> is identical to <code>0.0.0.0/0</code>• <code>1.2.3.4</code> is identical to <code>1.2.3.4/32</code>
external port	1. “any” 2. a single value 3. a range (e.g., <code>2000-3000</code>)	<ul style="list-style-type: none">• It specifies TCP/UDP port numbers. For ICMP packets, it is for the ICMP Type field.• The range is inclusive (i.e., <code>2000-3000</code> includes 2000, 2001, ..., and 3000)

DNS Rules

Format: <verdict> dns <domain name>

- verdict: “pass” or “drop”
- domain name: e.g., “bar.foo.com” (full domain name) or “*.gov” (wildcard domain name)
 - A full domain name is for **exact match**.
 - “*.gov” matches not only “fda.gov” but also “www.fda.gov”.
 - “*.foo.com” does not match “foo.com”
 - The asterisk (“*”) can only be used at the leftmost DNS label, alone.
 - Good syntax: “*”, “*.net”, “*.google.com”, ...
 - Bad syntax: “mail.*.com”, “*.foo.*.org”, “www*.salary.com”, ...

You apply DNS rules only for DNS query packets. More specifically, the packet should satisfy all of the following conditions to be considered for DNS rule matching.

- It is an outgoing UDP packet with destination port 53.
- It has exactly one DNS question entry.
 - There may be other non-empty sections (Answer, Authority, and Additional)
- The query type of the entry is either A or AAAA (QTYPE == 1 or QTYPE == 28), and
- The class of the entry is Internet (QCLASS == 1).

Example

Suppose that Starbucks wants to provide free WiFi, but with very restrictive rules as follows.

```
% allow “ping”, but no other types of ICMP packets
drop icmp any any
pass icmp any 0
pass icmp any 8

% allow DNS packets only to Google DNS servers
drop udp any any
pass udp 8.8.8.8 53
pass udp 8.8.4.4 53

% allow only HTTP(80), HTTPS(443), and Skype(1024-65535)
drop tcp any any
pass tcp any 80
pass tcp any 443
pass tcp any 1024-65535

% punish Italy (for not having Starbucks) and MIT (for the greedy /8 address block)
drop tcp it any
drop tcp 18.0.0.0/8 any

% ahem...
drop dns peets.com
drop dns *.peets.com
```

The rules implemented above do not allow any TCP/UDP/ICMP packets by default, but with some explicit exceptions.

Performance

Performance is one of the most important aspects of a firewall. Typically, commercial firewalls can process more than millions of packets per second. For this project, however, we will focus on its functionality and correctness, rather than performance. After all, the firewall is software-based, runs in a virtualized environment, and will be implemented in Python; it would be impossible to match the performance of commercial firewalls.

However, there will be minimum performance requirements. **If your firewall implementation underperforms by a factor of five when compared to our reference implementation with the same rule set, you may lose some points.** The performance will be measured in the number of processed packets, for a certain period of time (> 20 seconds, including startup time). Our reference implementation does not incorporate any sophisticated optimizations, so you do not need to worry much unless your implementation is very inefficient.

Notes

- Rule matching:
 - **If none of the rules match, just pass the packet.**
 - Thus you should always pass non-TCP/UDP/ICMP packets.
 - **If multiple rules match, apply the last one's verdict in the rules file.**
 - DNS request packets are inspected not only by DNS rules, but also by Protocol/IP/Port rules.
 - DNS rules may appear before Protocol/IP/Port rules.
 - For this project, we will not have more than 30 rules in the file, so linear scanning over the rules for every packet is perfectly fine.
 - The wildcard domain name matching described here does not conform to the standard (RFC 4592).
- Analyzing packets
 - `handle_packet()` will be called only for non-fragmented IPv4 packets.
 - Do not worry about packets with wrong TCP/IP checksum.
 - If a packet is so malformed that you cannot decode the packet for rule matching, simply drop the packet.
 - Also, if a DNS query packet (UDP with destination port 53) is too malformed to decode, you should drop the packet, even if there is no DNS rules.
 - If a DNS packet has non-question entries, you should consider the packet for DNS rules, as long as QDCOUNT==1 and the question entry conforms to the DNS rule matching conditions.
 - For example, `dig` will include one “Additional” entry in its DNS packets, but your firewall should apply DNS rules for those packets.
 - Your program should not crash.
 - Watch out for endianness. Most network protocols follow network order.
 - http://en.wikipedia.org/wiki/Endianness#Endianness_in_networking
- Parsing the rules file:
 - **All rules are case-insensitive, including domain names and country codes.** For example, “`tcp`”, “`TCP`”, and “`tCp`” must be all allowed.
 - You can assume that the syntax/format of the file is always correct. For example, we will not trick you with `128.50.132.20/24`, which has bad IP prefix syntax.
 - Ignore empty lines and comment lines (lines beginning with “`%`”)

GeoIP DB

In the `proj3` directory, we provide the database file `geoipdb.txt` that has geolocation-mapping information of the IP address space. You should use this file to implement the country-based blocking in your firewall. In the file, each line represents an IPv4 address range and its corresponding country code, in the following format.

Format

<start IP address> <end IP address> <2-letter ISO country code>

- start IP address, end IP address: These are dot-separated IPv4 addresses. The given IP address range is “inclusive” in that it includes both start and end IP addresses.
- 2-letter ISO country code:
 - http://en.wikipedia.org/wiki/ISO_3166-1_alpha-2
 - We use the same country code for the rules file (again, case-insensitive!)

The database file should be loaded into memory in `Firewall.__init__()`.

Notes

1. Note that we may use a different version of the database file for grading. However, the following two assumptions will always hold:
 - a. All IP address ranges given in the file do not overlap each other. Hence, no longest prefix matching will be required.
 - b. All IP address ranges are sorted in ascending order in the file.
2. Some IP addresses may not match any records in the database. It is normal, so don't worry.
3. Your lookup code for country matching should be “reasonably” fast, to meet the minimum performance requirements described above.
 - a. No linear search over the entire database for every packet!
 - b. Instead, consider using more efficient algorithms and data structures, such as binary search, radix tree, etc.
 - c. If loading of the database takes more than a few seconds, you are probably not on the right track.
4. For this project, we assume that the database is always correct and up-to-date.

The original database was retrieved from here: <http://dev.maxmind.com/geoip/legacy/csv/>

Bells and Whistles

You should consider doing some of the following three tasks, if you crave extra credits. Although they are not directly related to the core assignment, the firewall framework can be a great playground for fun experiments/features. While you are welcome to implement all of the them, you will only receive credit for up to two.

1. Packet Size Distribution

Measure the size of IP packets while surfing the Web in the bypass mode. We recommend measuring more than 50,000 of packets, for at least 30 minutes. Plot a bar graph with the packet size distribution in `bw1.pdf`. The x-axis represents 50-byte bins (1-50B, 51-100B, ..., 1451-1500B), and the y-axis is for the number of packets for each bin.

Describe the following in `bw1.txt`, in a few sentences for each.

- How long was the measurement period? How many packets did you measure in total?
- What did you mostly do during the period? (e.g., “watched kitty videos on Youtube”, “wasted my life at 9gag.com”, “Facebook stalked ex”, etc.)
- What does the distribution look like? Why does it show such a distribution?

Note: This task is measurement only, so your `firewall.py` file doesn’t need to include this feature. Only include the result (`bw1.pdf` and `bw1.txt`) in your submission tar file.

2. Loss Simulation

Implement the “`--loss <loss rate>`” option in `firewall.py`. The loss rate is always an integer. You can access command line options with the `config` dictionary given to the class constructor. With this option, your firewall should *randomly* drop packets (for both incoming and outgoing directions) at the specified loss probability. For example, “`sudo ./main.py --loss 5`” should drop 5% of packets, without considering firewall rules for those packets. Once you implement the option, do the following measurement.

Download a large file (> 10MB) with `wget` or `curl` with 0, 1, 3, 5, and 10% loss rates. For each loss rate, repeat the download at least for 5 times. Measure the TCP throughput (download speed), and plot the results (average and error bar for standard deviation) on a bar graph in `bw2.pdf`. In `bw2.txt`, summarize the result and explain the impact of loss rate on TCP throughput in one paragraph. Include both `bw2.pdf` and `bw2.txt` in your submission tar file.

3. Link Bandwidth Limit

Implement the “`--bwlimit <bandwidth>`” option in `firewall.py`. Again, you can access to command line options with the `config` dictionary given to the class constructor. With this option, your firewall should limit the maximum bandwidth of the link at the specified rate (for both incoming and outgoing directions). The rate is based on IP packet size, including the IP header.

The option argument is given in the form of “number[k|m]”. For example, “34567” (34,567 bits per second), “500k” (500,000 bits per second), or “2m” (2,000,000 bits per second). The number is always an integer.

To accomplish this task, implement two queues (one for each direction) to absorb transient traffic surge. Without the queues, TCP may perform poorly, due to its bursty behavior. Each queue should be able to store up to 500 packets. You will also need to use the timer functionality provided by the base code. Do not block the `handle_packet()` method with `time.sleep()`, a busy-wait loop, or any similar mechanisms.

References

The following documents provide the detailed protocol specifications to accomplish this project.

- IPv4 header:
 - <http://en.wikipedia.org/wiki/IPv4>
 - <http://tools.ietf.org/html/rfc791>
 - IP protocol numbers: <http://tools.ietf.org/html/rfc790>
- TCP header:
 - http://en.wikipedia.org/wiki/Transmission_Control_Protocol
 - <http://tools.ietf.org/html/rfc793>
- ICMP header:
 - <http://tools.ietf.org/html/rfc792>
- DNS packet format:
 - <http://tools.ietf.org/html/rfc1035>

You will (and should!) spend more time for testing than coding. We list some of the most helpful network testing tools below, all of which are preinstalled in the VM. There are a lot of online tutorials you can find on the Internet.

- **tcpdump / Wireshark**
 - **tcpdump** is a command-line packet sniffer that captures and display packets on networks. It also decodes packet headers for various protocols, which is very useful to verify the correctness of your own packet decoder.
 - **Wireshark** (formerly known as **Ethereal**) provides similar features with a graphical user interface.
 - They can also be used to check the behavior of your firewall. For example, if an outgoing packet is seen at both **int** and **ext**, it implies that the firewall successfully relayed the packet between those interfaces.
 - A short YouTube clip on Wireshark: <http://youtu.be/6X5TwvGXHP0>
- **nslookup / dig**
 - **nslookup** and **dig** are command-line tools for querying the Domain Name System.
 - You can use these tools to generate DNS query packets.
- **wget / curl**
 - **wget** and **curl** are handy tools to generate HTTP requests without using browsers.
- **nc**
 - **nc** (short for netcat) is a popular tool for generating TCP/UDP connections.
 - Refer to this page for more packet crafters: <http://sectools.org/tag/packet-crafters/>

We released an introductory document for these protocols/tools:

<http://www-inst.eecs.berkeley.edu/~ee122/fa13/projects/project3a/supp.pdf>

Logistics

Collaboration Policy

The project is designed to be solved independently, but you may work with at most one partner if you wish. Grading will remain the same whether you choose to work alone or in partners; both partners will receive the same grade regardless of the distribution of work between the two partners (so choose a partner wisely!).

By the nature of this project, it could be hard to “split” the work between teammates. Instead, consider pair programming (http://en.wikipedia.org/wiki/Pair_programming), if you choose to work in partners.

Submission Instructions

Turn in a **.tar** file with both you and your partner’s last names in the file name (For example, `project3a-han-shrikhande.tar` or `project3-panda.tar`) on an instructional machine. Your tar file should include:

- `firewall.py`: remember, this file should include all your code, including B&W 2 and 3.
- `readme.txt`
 - Your (and your partner’s) full name and login name
 - (optional) Any comments/concerns on the project. This will not affect grading. We will collect your opinions anonymously to design future projects better.
- (optional, for B&W 1) `bw1.pdf` and `bw1.txt`
- (optional, for B&W 2) `bw2.pdf` and `bw2.txt`

```
$ tar -cf project3a-han-shrikhande.tar firewall.py readme.txt [bonus files ...]
$ submit project3a
```

Regrade & Late Policy

Your regrade request must be made within seven days after score release. Submit your new tar file on an instructional machine, then email sangjin@eecs.berkeley.edu.

- $\text{final score} = \max\{\text{old score}, \text{new score} * (\text{regrading penalty})\} * (\text{your original late penalty})$
 - $\text{regrading penalty} = 0.9 - (\text{code difference in bytes}) * 0.001$
 - We will provide a script that measures code difference.
 - Of course, regrading penalty doesn’t apply if it turns out be auto-grader’s fault.
- No regades for bells and whistles.
- You have only one opportunity for regrade request.

The late policy is simple; no slip dates. If submission is late, it is penalized as follows:

- < 24 hours late, you lose 10%
- < 48 hours, 20%
- < 72 hours, 40%
- More than three days late, you can no longer hand-in the assignment.

Cheating

Simply put, DO NOT EVER TRY IT. We will do our best to protect our students from losing the value of their honestly earned grades due to cheaters. We may perform manual inspection and use automated tools (do not underestimate them) to detect potential plagiarism over all submissions.

Refer to [the course webpage](#) for more information. If you are not sure what may constitute cheating, consult the instructor or GSIs. Assignments suspected of cheating or forgery will be handled according to the Student Code of Conduct². Apparently 23% of academic misconduct cases at a certain junior university are in Computer Science³, but we expect you all to uphold high academic integrity and pride in doing your own work.

² <http://sa.berkeley.edu/code-of-conduct>

³ http://www.pcworld.com/article/194486/Computer_Science_Students_Cheating.html

DO's and DON'Ts

DO's

- It is okay to use external libraries or applications to **test** your firewall.
- You can modify not only `firewall.py` but also other source code files, but only for debugging purposes.
 - Do remember that you only submit the `firewall.py` file, and it should work with the original code of other files.
- Backup is always a good idea.
- Feel free to surf the Web for **general ideas**, such as Python language/library references, network protocol specifications, testing tool usages, and debugging tips.
- You can discuss with anyone about algorithms, approaches, and concepts without details, but stay away from your computer/code while doing so.
- We encourage you to share test strategies with your fellow students on Piazza, but only at the high level (e.g., no test code).
- We recommend using Firefox in the VM, rather than installing other web browsers. The Firefox installed in the VM was specially configured to suppress some seemingly strange behaviors.

DON'Ts

- Do not create extra threads or processes. Use the provided timer functionality instead.
- Do not use any libraries other than Python 2.7 standard modules to implement the firewall.
- Do not alter the network configurations of the VM. Do not install “Network Manager” package. The VM relies on manual/delicate configurations to provide the firewall functionalities. You may have to reinstall the VM if some configuration gets broken.
- **All parts of the solution code must be your own; copying someone else's code snippet is strictly prohibited.**
- Do not share your code with anyone, other than your project partner.
- Do not post any project-specific questions on Internet forums other than Piazza.
- The project is led by Sangjin Han (main), Steve Wang, Kaifei Chen, and Aurojit Panda. Avoid asking other GSIs project-specific questions.